

Effective Programming Practices for Economists

# Scientific Computing

Writing fast code with numba

Janoš Gabler and Hans-Martin von Gaudecker

# Why numba can be fast

- Python code is just-in-time compiled to machine code
- It builds on top of numpy
  - Efficient storage of data in arrays
  - All dtypes known at compile-time
- Uses same technology as Julia under the hood

# Implications

- Loops can be super fast
- Avoids some of the drawbacks of vectorized code
- However: Not guaranteed to be faster than numpy

# Numba is picky

- Need to write out all loops
- Works best if all your data is in scalars or arrays
  - No dicts, lists or NamedTuples
- Inside compiled functions you cannot call uncompiled functions
  - Cannot call most libraries
  - Sometimes need to re-implement algorithms in numba compatible ways

Therefore, numba is intended to speed up your bottlenecks, not to compile your entire program

# Naively jitting the example

```
from numba import njit

@njit
def array_cobb_douglas(factors, weights, a):
    out = np.empty(len(factors))
    for i in range(len(factors)):
        out[i] = _cobb_douglas(factors[i], weights, a)
    return out
```

```
@njit
def _cobb_douglas(factors, weights, a):
    return a * np.prod(factors**weights)
```

# (inputs as before)

```
%timeit array_cobb_douglas(factors, weights, a)
```

1.22 ms ± 39.6 µs per loop

- The function is numba compatible but will not be optimal because not all loops are written out
- It still becomes much faster than before (~20× speedup)

# Full loops

```
@njit
def numba_array_cobb_douglas(factors, weights, a):
    out = np.empty(len(factors))
    for i in range(len(factors)):
        out_i = a
        for j in range(len(weights)):
            out_i *= factors[i, j]**weights[j]
        out[i] = out_i
    return out
```

```
%timeit numba_array_cobb_douglas(factors, weights, a)
```

602  $\mu$ s  $\pm$  3.62  $\mu$ s per loop

- Writing out loops yields gets us another  $\sim 2\times$  speedup ( $\sim 40\times$  over original)
- Here, still slower than vectorized numpy but this is not always the case!

