Effective Programming Practices for Economists

Data management with pandas

Data types

Janoś Gabler and Hans-Martin von Gaudecker

Overview

- Why different data types?
- Converting to efficient dtypes
- Overview of numeric dtypes
- String vs. Categorical
- Working with strings and categoricals

The need for different data types

Consider the gapminder data

	country	continent	year	life_exp
0	Cuba	Americas	2002	77.16
1	Cuba	Americas	2007	78.27
2	Spain	Europe	2002	79.78
3	Spain	Europe	2007	80.94

```
>>> df.dtypes

country       string[pyarrow_numpy]
continent     string[pyarrow_numpy]
year                          int64
life_exp                    float64
dtype: object
```

- Each column has a dtype
- Enables efficient storage and fast computation
- Dtypes are not always set optimally after loading data

Benefits of good type representation

- Fast calculations in a low level language
- Access to operations that are only relevant for some types
- Memory efficiency

Converting to efficient dtypes

```
>>> better_dtypes = {
...     "country": pd.CategoricalDtype(),
...     "continent": pd.CategoricalDtype(),
...     "year": pd.UInt16Dtype(),
...     "life_exp": pd.Float64Dtype(),
... }

>>> df = df.astype(better_dtypes)
>>> df.dtypes

country       category
continent     category
year             UInt16
life_exp        Float64
dtype: object
```

- Depending on how you load your data, the dtypes are not set optimally
- If so, you can create a dictionary that maps columns to the dtypes you want

Overview of numeric dtypes

Type	Properties
`pd.Int8Dtype()`	Byte (-128 to 127)
`pd.Int16Dtype()`	Integer (-32768 to 32767)
`pd.Int32Dtype()`	Integer (-2147483648 to 2147483647)
`pd.Int64Dtype()`	Integer (-9223372036854775808 to 9223372036854775807)
`pd.UInt8Dtype()`	Unsigned integer (0 to 255)
`pd.UInt16Dtype()`	Unsigned integer (0 to 65535)
`pd.UInt32Dtype()`	Unsigned integer (0 to 4294967295)
`pd.UInt64Dtype()`	Unsigned integer (0 to 18446744073709551615)
`pd.Float64Dtype()`	Double precision float

String vs. Categorical

- `pd.CategoricalDtype()` is for data that takes values in a fixed and relatively small set of categories
  - Internally stored as small integers
  - Very fast relabeling or resorting of categories
- `pd.StringDtype()` is for actual text data
  - Internally stored as `pyarrow` array
  - Fast string functions similar to methods of Python strings

Working with strings

```
>>> sr = pd.Series(["Guido", "Tim", "Raymond"])

>>> sr.str.lower()

0      guido
1        tim
2    raymond
dtype: string


>>> sr.str.replace("i", "iii")

0    Guiiido
1      Tiiim
2    Raymond
dtype: string
```

- The `.str` accessor provides access to the string methods

- Vectorized and fast implementations!

- Other examples:

  - `sr.str.len`

  - `sr.str.contains`

  - ...

- See this tutorial for more string methods

Working with categoricals

```
>>> cat_type = pd.CategoricalDtype(
...     categories=["low", "middle", "high"],
...     ordered=True,
... )

>>> sr = pd.Series(
...     ["low", "high", "high"],
...     dtype=cat_type,
... )

>>> sr
0      low
1     high
2     high
dtype: category
Categories (3, string): [low < middle < high]
```

- Categories are defined independent of data
  - Protection against invalid categories
  - Good for visualization!
- `sr.cat` accessor provides access to methods
- See this tutorial for more methods